



# Optimal spike-based signal representation on a neuromorphic chip



# Outline

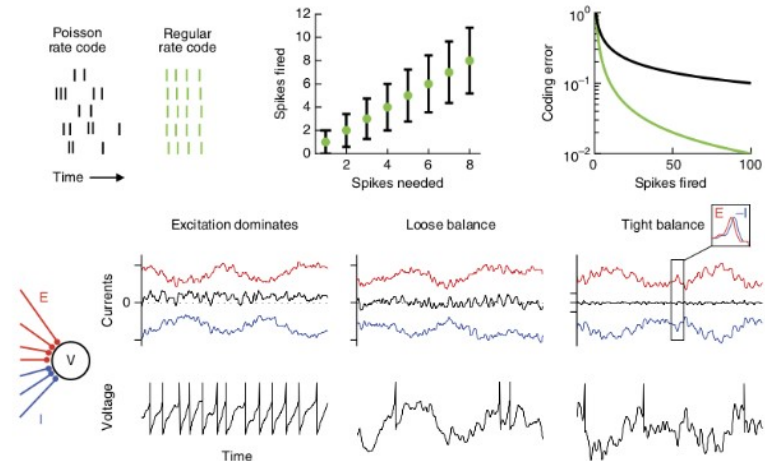
- Introduction & Problem statement
- Goals of the thesis
- The DYNAP-SE
- Learning in-the-loop
- Experiments & Discussion

# Signal representation is a fundamental problem

- Unclear how population of neurons represent a signal *effectively* with usable precision
- Coding error using Poisson rate codes scales with  $\frac{1}{\sqrt{M}}$
- Representing multiple signals requires many spikes  $\rightarrow$  waste of energy (caused by irregularity of Poisson codes)
- Goal: Learning rule driving network in state where
  - Coding error scales with  $\frac{1}{M}$
  - Code is *robust* to noise
  - Code is *energy efficient*

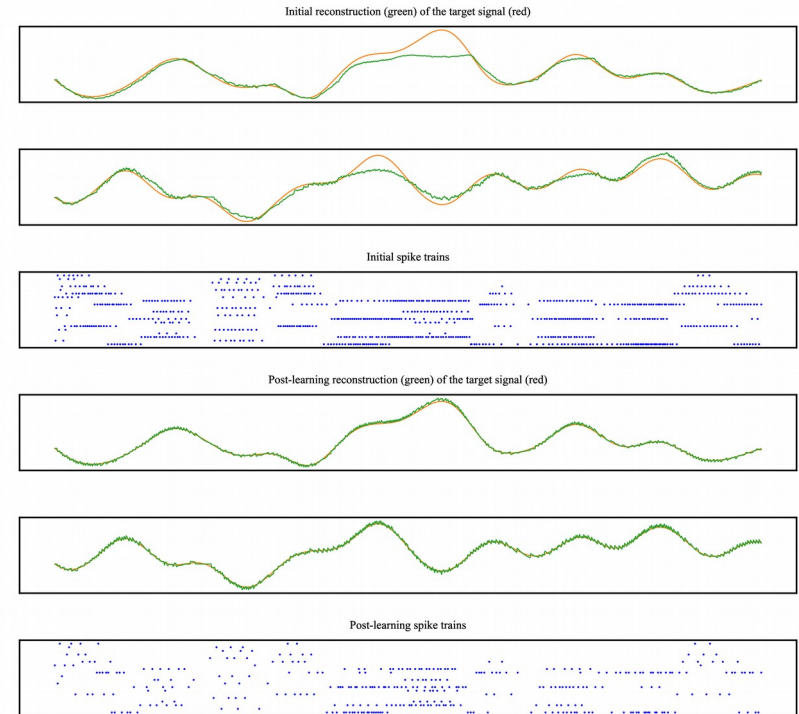
# Possible solution: Tightly balanced networks

- Balance achieved by matching excitatory with inhibitory input (E/I-balance)
- Balanced networks produce Poisson-like codes
- Tight balance: Small fluctuations are matched (e.g. excitation followed by inhibition)



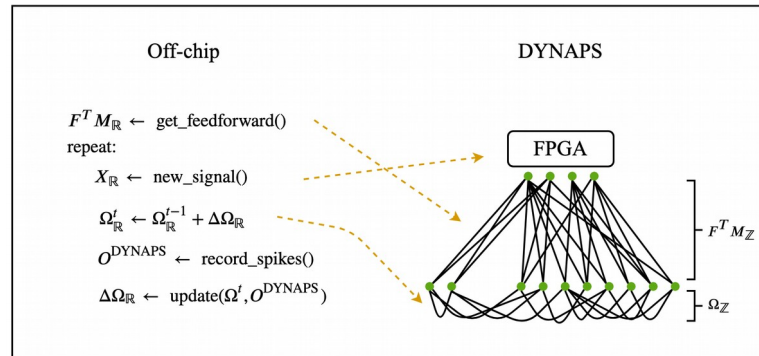
# Learning to represent signals spike by spike

- Key insights/contributions:
  - Neural circuits behave accurately and every spike “counts”
  - Derivation of local online learning rule:
    - Condition for spike: Neuron reduces reconstruction error (*error-driven coding*)  
→ Derivation of optimal weights
    - Minimizing voltage deviations → reaching optimal weights → low reconstruction error




# Goals of the thesis

- Implement this learning rule on neuromorphic chip (DYNAP-SE) in-the-loop
- Lay foundation to implement directly on-chip
- Identify possible weaknesses and suggest improvements



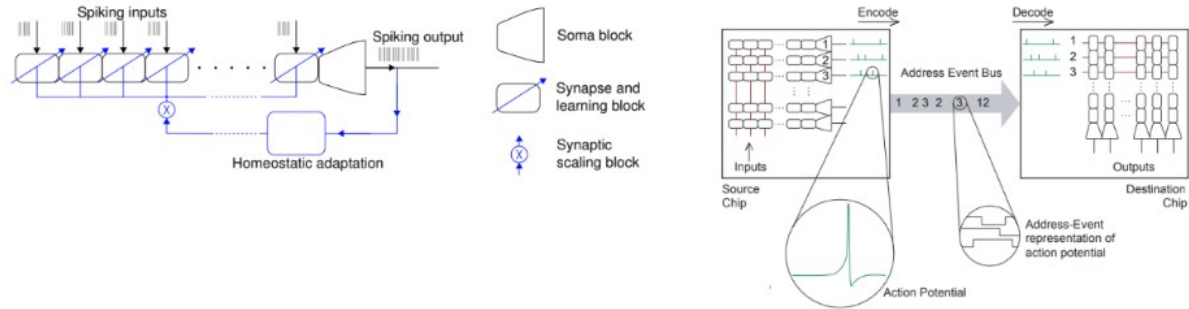
# DYNAP-SE

- Mixed analog/digital circuits
- Completely asynchronous (no central clock)
- Ultra low power consumption e.g. 276uA @ 1.3V with all neurons firing at 30Hz (CPU uses 30W @14V = 2.1A)
- Fan-in: 64 | Fan-out: 4064
- 256 silicon neurons/core, 64 synapses/neuron
- Neuron model: Adaptive-Exponential I&F

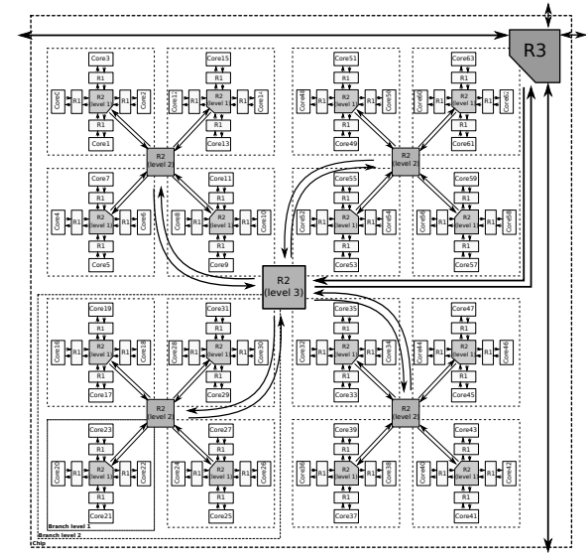




# DYNAP-SE - Routing




- **Address-Event-Representation**
- Combination of hierarchical and 2D-mesh routing
- Why hierarchical? Smaller address space when assuming locality → save memory
- Spike times encoded in Address Event times





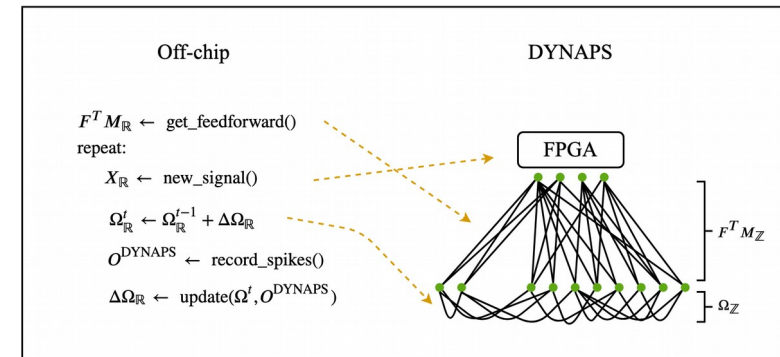
# Demo



C0C0	C0C1	C0C2	C0C3	C1C0	C1C1	C1C2	C1C3	C2C0
2400000000 1568628	2400000000 0	2400000000 1568628	2400000000 0	2400000000 35294120	2400000000 0	2400000000 0	2400000000 58847	
IF_AHTAU_N	IF_AHW_P	IF_BUF_P	IF_CASC_N	IF_DC_P	IF_NMDSA_N	IF_RFR_N		
2400000000 29	2400000000 588	2400000000 282353	2400000000 235	2400000000 0	2400000000 0	2400000000 235		
IF_TAU1_N	IF_TAU2_N	IF_THR_N	NPDPIE_TAU_F_P	NPDPIE_TAU_S_P	NPDPII_TAU_F_P	NPDPII_TAU_S_P		
2400000000 0	2400000000 0	2400000000 0	2400000000 0	2400000000 2078431	2400000000 200000			
PS_WEIGHT_EXC_F_NPS_WEIGHT_EXC_S_N	PS_WEIGHT_INH_F_N	PS_WEIGHT_INH_S_N	PULSE_PWLK_P	R2R_P				


# Learning in-the-loop

- Only train recurrent connections
- Mismatches between simulation and in-the-loop implementation:
  - Input must be in spike-form
  - Weight updates must be batched
  - Weight matrices must be discrete
  - Cannot measure voltage in software → Need to estimate from spikes and input
  - Update only neuron with highest voltage above threshold



# Learning in-the-loop - Input

- Use delta modulation to transform continuous signal to spikes
- Could simply integrate signal using  $\hat{x}_t = (1 - \lambda)\hat{x}_{t-1} + \mu(o_t^{up} - o_t^{down})$
- General form for 2 inputs :  $\hat{x}_t = (1 - \lambda)\hat{x}_{t-1} + \mu M o_{t-1}$  with  $M = \begin{bmatrix} 1 & -1 & 0 & 0 \\ 0 & 0 & 1 & -1 \end{bmatrix}$
- → Feed-forward matrix becomes  $F^T M$  and integration is performed by synapses



# Learning in-the-loop - Updates

- In the simulation: Update rate and weights *only* for neuron highest above the threshold
- To align with the DYNAP-SE, check if *assumption* of highest neuron above threshold matches DYNAP-SE recordings within some time-frame

```


(max, k) ← arg max(Vt - thresh - 0.01 · N(0, 1))
if max ≥ 0 then // Neuron k spiked
    | Ωk ← Ωk - εΩ · (β · (Vt + μ · rt-1) + Ωk + μ · Ik)
    | rt ← rt-1 + Ik
else
    | rt ← rt-1
end
rt ← (1 - λdt) · rt
  
```

```

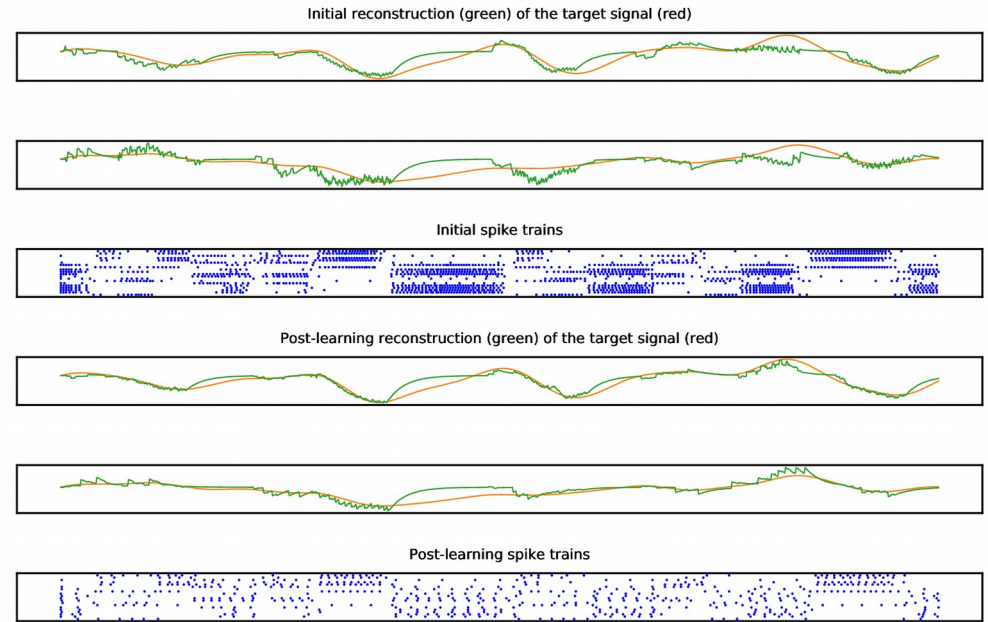
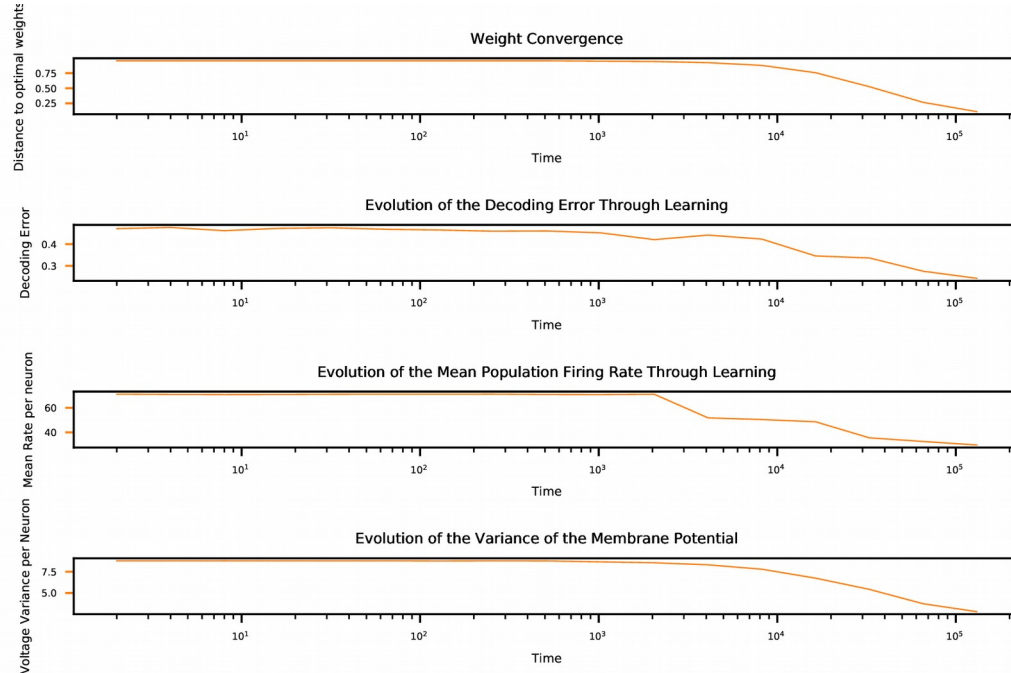
(max, k) ← arg max(Vt - thresh - 0.01 · N(0, 1))
if max ≥ 0 ∧ has_spike(Ok, ±ΔTDYNAPSE) then
    | ΔΩk ← ΔΩk + εΩ · (β · (Vt + μ · rt-1) + Ωk + μ · I)
    | otaligned ← Ik
    | rt ← rt-1 + Ik
else
    | otaligned ← 0
    | rt ← rt-1
end
rt ← (1 - λdt) · rt
  
```

# Experiments & Discussion

- Used 20 Neurons, trained for 140 iterations, 2 input signals
- Used Reinforcement Learning based approach to adapt learning rate
  - Reward = Reconstruction error
  - Action = Scale learning rate
- Intuition: Rewards for good signal representation should keep weights in current regime



# Experiments & Discussion



## Experiments & Discussion

- Very close to simulation → constant excitement would render condition always true and produce same results as simulation
- Update rule from the paper utilizes voltage *before* spike propagated, but can rewrite to use voltage after

$$n \leftarrow \arg \max (\mathbf{V} - \mathbf{T}) - \xi_T(\tau))$$

if  $V_n > T_n$  then

$$o_n(\tau) = 1$$

$$\mathbf{F}_n(\tau) = \mathbf{F}_n(\tau - 1) + \epsilon_F(\alpha \mathbf{x}(\tau - 1) - \mathbf{F}_n(\tau - 1))$$

$$\Omega_n(\tau) = \Omega_n(\tau - 1) - \epsilon_\Omega(\beta(\mathbf{V}(\tau - 1) + \mu \mathbf{r}(\tau - 1)) + \Omega_n(\tau - 1))$$

- Last problem: Learning fails when discretizing weights  
Derive for discrete weights?
- Goal: Resolve issues and implement in circuit on next chip-iteration